

Undirected Graph

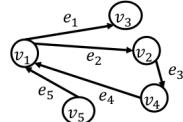
- An undirected graph is a pair of (V, E) where:
 - V is a set of elements, each of which called a node
 - E is a set of unordered pairs $\{u, v\}$ such that u and v are nodes

Directed Graph

- An directed graph is a pair of (V, E) where:
 - V is a set of elements, each of which called a node
 - E is a set of ordered pairs $\{u, v\}$ where u and v are nodes, we say there is a directed edge from u to v .
- A directed edge (u, v) is said to be an outgoing edge of u , and incoming edge of v . Accordingly, v is an out-neighbor of u , and u is in-neighbor of v .

Definitions of Graph

- Let $G = (V, E)$ be a graph. A path in G is a sequence of nodes (v_1, v_2, \dots, v_k) such that
 - For every $i \in [1, k - 1]$, there is an edge between v_i and v_{i+1} .
- A cycle in G is a trail in which the only repeated vertices are the first and last vertices.
- Example:
 - Cycle: (v_1, v_2, v_4, v_1) ; Path: (v_5, v_1, v_2, v_4)
- In an undirected graph, the degree of vertex u is the number of edges of u
- In a directed graph, the out-degree of a vertex u is the number of outgoing edges of u , and its in-degree is the number of its incoming edges



在无向图中,

$$\sum_{u \in V} d(u) = 2|E|$$

In degree : $d^+(v_i) = 2$

out degree : $d^-(v_i) = 2$

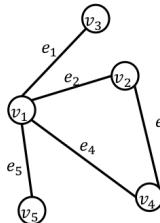
在有向图中,

$$\sum_{u \in V} d^+(u) = |E|$$

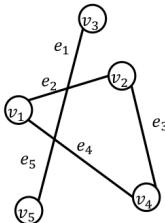
$$\sum_{u \in V} d^-(u) = |E|$$

Connected Graph (连通图)

- An undirected graph $G = (V, E)$ is connected if, for any two distinct vertices u and v , G has a path from u to v .



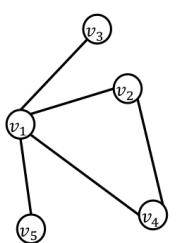
connected



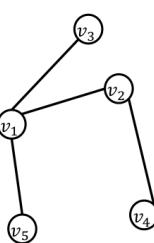
not connected

Graph. Tree. Forest

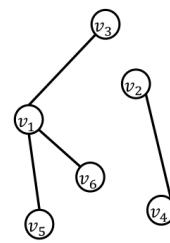
- A tree is a connected undirected graph contains no cycles.
- Forest is a set of disjoint trees.



Graph, not tree



Graph, tree

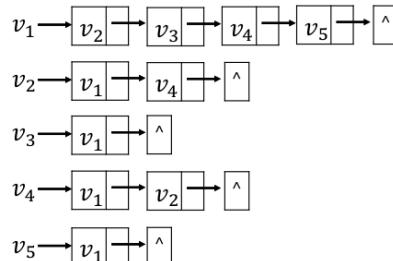
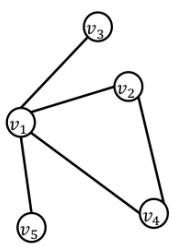


Graph, forest

Adjacency List: 邻接表

Undirected Graph

- Each vertex $u \in V$ is associated with a linked list that enumerates all the vertices that are connected to u .

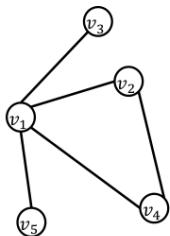


Space = $O(|V| + |E|)$

Adjacency Matrix: 邻接矩阵

Undirected Graph

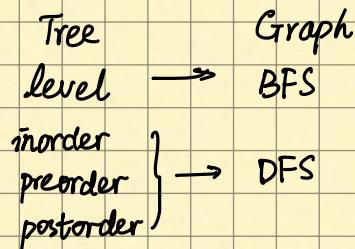
- A $|V| * |V|$ matrix A where $A[u,v] = 1$ if $(u,v) \in E$, or 0 otherwise



	v_1	v_2	v_3	v_4	v_5
v_1	0	1	1	1	1
v_2	1	0	0	1	0
v_3	1	0	0	0	0
v_4	1	1	0	0	0
v_5	1	0	0	0	0

- A must be symmetric
- Space = $O(|V|^2)$

Graph Traversal

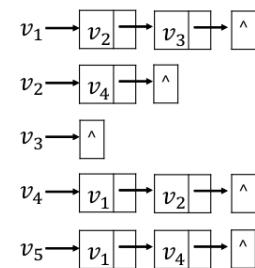
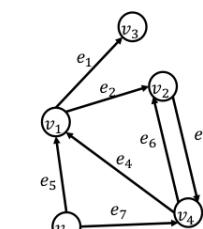


Shortest Path

- Let $G = (V, E)$ be a directed graph. A path in G is a sequence of nodes (v_1, v_2, \dots, v_k) such that
 - For every $i \in [1, k - 1]$, there is an edge between v_i and v_{i+1} .
 - E.g., $(v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v_k)$
 - Sometimes, we also denote the path as $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k$
- The path is said to be from v_1 to v_k , the length of the path is $k - 1$.
- Given two vertices $u, v \in V$, a shortest path from u to v is a path from u to v that has the minimum length among all the paths from u to v .
- If there is no path from u to v , then v is said to be unreachable from u .

Directed Graph

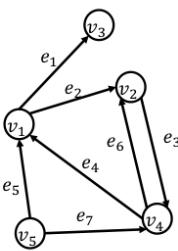
- Each vertex $u \in V$ is associated with a linked list that enumerates all the vertices $v \in V$ that there is an edge from u to v .



Space = $O(|V| + |E|)$

Directed Graph

- Defined in the same way as in the undirected graph



	v_1	v_2	v_3	v_4	v_5
v_1	0	1	1	0	0
v_2	0	0	0	1	0
v_3	0	0	0	0	0
v_4	1	1	0	0	0
v_5	1	0	0	0	0

- A may not be symmetric.
- Space = $O(|V|^2)$

Breadth First Search (BFS)

- At the beginning, color all vertices in graph white. And create an empty BFS tree T.
- Create a queue Q. Insert the source vertex s into Q, and color it yellow (which means "in the queue")
- Make s the root of T.
- Repeat the following until Q is empty
 - De-queue from Q the first vertex v
 - For every out-neighbor u of v that is still white
 - Enqueue u into Q, and color u yellow
 - Make u a child of v in the BFS tree T.
 - Color v red (meaning v is visited)

Single Source Shortest Path (SSSP)

从一个源节点出发，到其他节点的最短路径。

BFS树中源a到任意节点x的路径即最短路径

初始化 graph G 全部节点为白色。创建一个空的BFS树 T。

创建一个队列 Q，往 Q 中加入源节点 s，并将其变成黄色。

令 s 为 T 的根节点。

while (!Q.isEmpty())

v = Q.dequeue() v ← red

for (v's neighbour u)

if (u is still white)

Q.enqueue(u) u ← yellow

T(u is v's child)

endif

endfor

endwhile

事实上黄/红相当于一种颜色

Complexity Analysis

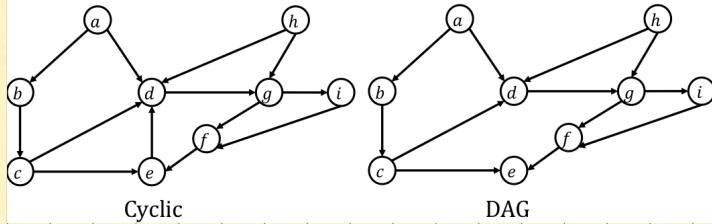
- When a vertex v is dequeued, we spend $O(1 + d^+(v))$ time processing it, where $d^+(v)$ is the out-degree of v.
- Clearly, every vertex enters the queue at most once.
- The total running time of BFS is therefore:

$$O\left(\sum_{v \in V} (1 + d^+(v))\right) = O(|V| + |E|)$$

$O(1 + d^+(v)) =$ 从一个节点 + 该节点的 in degree

Directed Acyclic/Cyclic Graph

- ◆ If a directed graph contains no cycles, we say that it is a directed acyclic graph (DAG). Otherwise, G is Cyclic.
- ◆ DAG is extremely important concept in Computer Science, e.g., spark, tensorflow
- ◆ Example



有向无环图 DAG

Depth First Search (DFS)

- ◆ At the beginning, color all vertices in the graph white, and create an empty DFS tree T.
- ◆ Create a stack S. Pick an arbitrary vertex v. Push v into S, and color it yellow (which means "in the stack")
 - ◆ What is the difference between BFS and DFS underlying data structure?
 - ◆ BFS → Queue, DFS → Stack
- ◆ Make v the root of T
- ◆ Repeat the following until S is empty
 - ◆ Let v be the vertex that currently tops the stack S (do not remove v from S)
 - ◆ Does v still have a white out-neighbor
 - ◆ 2.1 If yes: let it be u.
 - ◆ Push u into S, and color u yellow
 - ◆ Make u a child of v in the DFS-tree T
 - ◆ 2.2 If no, pop v from S, and color v red (meaning v is visited)
 - ◆ If there are still white vertices, repeat the above by restarting from an arbitrary white vertex v', creating a new DFS tree rooted at v'.

初始给定graph G, 全部节点白色, 创建空DFS树T.

创建栈S. 取出一节点v, 加入S, 并改成黄色.

令v为T的根节点.

while (!S.isEmpty())

v = S.top()

for (v's neighbour u)

if (u is still white)

S.push(u) u ← yellow

T.u is v's child)

endif

endfor

S.pop() (将节点v pop出去) v ← red

endwhile

若图中仍有白色节点, 则任取一个加入S中继续DFS.

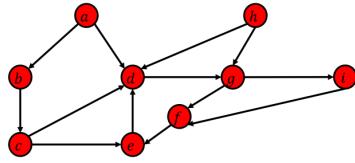
DFS最终也会生成一棵tree(或者forest)

Complexity Analysis

- ◆ DFS can be implemented efficiently as follows.
 - ◆ Store G in the adjacency list format
 - ◆ For every vertex v, remember the out-neighbor to explore next
 - ◆ $O(|V|+|E|)$ stack operations
 - ◆ Use an array to remember the colors of all vertices
- ◆ Hence, the total running time is $O(|V|+|E|)$.

Edge classification

- Suppose we have already built a DFS-forest T.
- Let (u,v) be an edge in G (remember that the edge is directed from u to v). It can be classified into:
 - Forward edge: u is a proper ancestor of v in a DFS-tree of T.
 - Backward edge: u is a descendant of v in a DFS-tree of T.
 - Cross edge: if neither of the above applies.



DFS Forest

a
↓
b
↓
c
↓
d
↓
g
↓
f
↓
e

- $I(a) = [1, 16]$, $I(b) = [2, 15]$, $I(c) = [3, 14]$
- $I(d) = [4, 13]$, $I(g) = [5, 12]$, $I(f) = [6, 9]$
- $I(e) = [7, 8]$, $I(i) = [10, 11]$, $I(h) = [17, 18]$

- Forward edge:
 - $(a,b), (a,d), (b,c), (c,d), (c,e), (d,g), (g,f), (g,h), (f,e)$
- Backward edge: (e,d)
- Cross edge: $(i,f), (h,d), (h,g)$

$O(1)$ 确定 $edge(u,v)$ 的类型

- Maintain a counter c , which is initially 0. Every time a push or pop is performed on the stack, we increment c by 1.
- For every vertex v , define:
 - Its discovery time $d-tm(v)$ to be the value of c right after v is pushed into the stack
 - Its finish time $f-tm(v)$ to be the value of c right after v is popped from the stack
 - Define $I(v) = [d-tm(v), f-tm(v)]$
- It is straight forward to obtain $I(v)$ for all $v \in V$ by paying $O(|V|)$ extra time on top of DFS's running time.

设立一个计数器 $c=0$. 每当有节点 a 被 push 或 pop 出 stack 时, $c++$.

对每个节点 v , 有:

$d-tm(v)$: v 被 push 入 stack 时的 c 值.

$f-tm(v)$: v 被 pop 出 stack 时的 c 值

定义 $I(v) = [d-tm(v) f-tm(v)]$

Theorems

- Parenthesis Theorem:** all the following are true:
 - If u is a proper ancestor of v in DFS-tree of T, then $I(u)$ contains $I(v)$.
 - If u is a proper descendant of v in DFS-tree of T, then $I(u)$ is contained in $I(v)$.
 - Otherwise, $I(u)$ and $I(v)$ are disjoint.
- Proof: Follows directly from the first-in-last-out property of the stack.
- Cycle Theorem:** let T be an arbitrary DFS-forest. G contains a cycle if and only if there is a backward edge with respect to T.
- Proof: will left as exercise.

u 是 v 的真祖先 $\Leftrightarrow I(v) \in I(u)$

u 是 v 的真后代 $\Leftrightarrow I(u) \in I(v)$

否则, $I(u)$ 与 $I(v)$ 不相交.

G 有环 \Leftrightarrow T 中有 backward edge.

利用 cycle Theorem.

- Equipped with the cycle theorem, we know that we can detect whether G has a cycle easily after having obtained a DFS-forest T:
 - For every edge (u,v) , determine whether it is a backward edge in $O(1)$ time.
- If no backward edges are found, decide G to be a DAG; otherwise, G has at least a cycle.
- Only $O(|E|)$ extra time is needed
- We now conclude that the cycle detection problem can be solved in $O(|V|+|E|)$ time.

Topological Sort

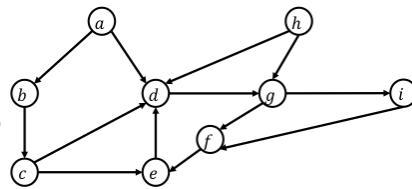
- Let $G=(V,E)$ be a directed acyclic graph (DAG).
- A topological order of G is an ordering of the vertices in V such that, for any edge (u,v) , it must hold that u precedes v in the ordering.

Example: two possible topological orders:

- $h, a, b, c, d, g, i, f, e$
- $a, g, b, c, d, g, i, f, e$

- $a, h, d, b, c, g, i, f, e$

is not topological order, because of edge (c,d) .

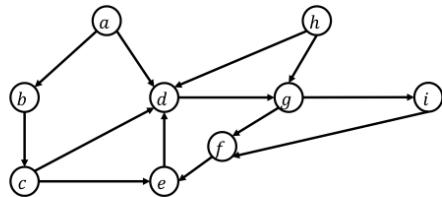


- Let $G=(V,E)$ be a directed acyclic graph (DAG). The goal of topological sort is to produce a topological order of G .

Topological Sort Algorithm

- Create an empty list L
- Run DFS on G , whenever a vertex v turns red (i.e., it is popped from the stack), append it to L .
- Output the reverse order of L

- The total running time is clearly $O(|V|+|E|)$



- Suppose we run DFS starting from a . The following is one possible order by which the vertices turn red:

- $e, f, i, g, d, c, b, a, h$

- Therefore, we output $h, a, b, c, d, g, i, f, e$ as a topological order.

- Suppose we run DFS starting from d , then restarting from h , then from a . The following is one possible order by which the vertices turn red:

- $e, f, i, g, d, h, c, b, a$

- Therefore, we output $a, b, c, h, d, g, i, f, e$ as a topological order.

拓扑序是对任意边 (u,v) , u 都是 v 的前趋。

拓扑排序的第一个节点必须没有进边，最后一个节点必须没有出边。

拓扑序不只一个。

(因为源节点不一定一样)

每一个 DAG 都有 topological order

创建空列表 L

对 G 跑 DFS, 当有节点 v 出栈时, 加在末尾
最终 L 的逆序即拓扑序。

Correctness Analysis

- We now prove that the algorithm is correct.

- Proof. Take any edge (u,v) . We will show that u turns red after v , which will complete the proof.

- Consider the moment when u enters the stack. We argue that currently v cannot be in the stack. Suppose that v was in the stack. As there must be a path chaining up all the vertices in the stack bottom up, we know that there is a path from v to u . Then, adding the edge (u,v) forms a cycle, contradicting the fact that G is a DAG.

- v is red at this moment then obviously u will turn red after v .

- v is white: then by the white path theorem of DFS, we know that v will become a proper descendant of u in the DFS-forest. Therefore, u will turn red after v .

- Every DAG has a topological order!

Shortest Path (最短路径问题)

Single source shortest path (SSSP)

- BFS algorithm
- All the edges have the same weight

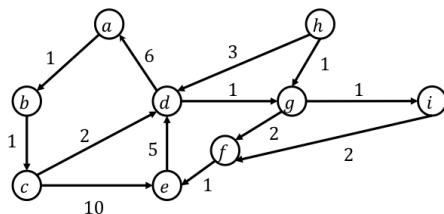
SSSP with arbitrary positive path (SP)

Weight graph

- Let $G=(V,E)$ be a directed graph. Let w be a function that maps each edge in E to a positive integer value. Specifically, for each $e \in E$, $w(e)$ is a positive integer value, which we call the weight of e .
- A directed weighted graph is defined as the pair (G,w) .
- Consider a directed weighted graph defined by a directed graph $G=(V,E)$ and function w .
- Consider a path in G : $(v_1, v_2), (v_2, v_3), \dots, (v_l, v_{l+1})$, for some integer $l \geq 1$. We define the length of the path as: $\sum_{i=1}^l w(v_i, v_{i+1})$.
- Recall that we may also denote the path as: $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_{l+1}$.
- Give two vertices $u, v \in V$, a shortest path from u to v is a path from u to v that has the minimum length among all the paths from u to v .
- If v is unreachable from u , then the shortest path distance from u to v is ∞ .

SSSP with Positive Weights

- Let (G,w) with $G=(V,E)$ be a directed weighted graph, where w maps every edge of E to a positive value.
- Give a vertex s in V , the goal of the SSSP problem is to find, for every other vertex $t \in V \setminus \{s\}$, a shortest path from s to t , unless t is unreachable from s .
- A subsequence property
 - Lemma: if $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_{l+1}$ is a shortest path from v_1 to v_{l+1} , then for every i, j satisfying $1 \leq i \leq j \leq l + 1$, $v_i \rightarrow v_{i+1} \rightarrow \dots \rightarrow v_j$ is shortest path from v_i to v_j .
 - Proof: suppose that this is not true, then we can find a shorter path from v_i to v_j . Using that path to replace the original path from v_1 to v_{l+1} , which contradicts the fact that $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_{l+1}$ is a shortest path.



- The path $c \rightarrow e$ has length 10
- The path $c \rightarrow d \rightarrow g \rightarrow f \rightarrow e$ has length 6
- The second path is the shortest path from c to e
- We know that any subsequence of this path is also a shortest path. For example, $c \rightarrow d \rightarrow g \rightarrow f$ must be a shortest path from c to f .

Dijkstra's Algorithm

edge relaxation

The edge relaxation idea

- For every vertex $v \in V$, we will maintain a value $\text{dist}(v)$ that represents the length of the shortest path from s to v found so far.
- At the end of the algorithm, we will ensure that every $\text{dist}(v)$ equal to the precise shortest path from s to v .
- A core operation in our algorithm is called edge relaxation. Given an edge (u,v) , we relax it as follows:
 - If $\text{dist}(v) < \text{dist}(u) + w(u,v)$, do nothing
 - Otherwise, reduce $\text{dist}(v)$ to $\text{dist}(u) + w(u,v)$

71

对每个节点 v , 都有到节点 s 的最短距离 $\text{dist}(v)$

对某条边 $\text{edge}(u,v)$, 若 $\text{dist}(v) > \text{dist}(u) + w(u,v)$ 就令 $\text{dist}(v)$ 等于 $\text{dist}(u) + w(u,v)$

初始化时将所有节点 v 的父亲 $\text{parent}(v)$ 设为 null.

对源点 s 设 $\text{dist}(s)=0$, 其他 $\text{dist}(v)=\infty$.

设 $S = V$.

while (!S.isEmpty())

$u = \text{smallest dist}(u)$ in S

 for out edge (u,v) of u

 if $\text{dist}(v) > \text{dist}(u) + w(u,v)$

$\text{dist}(v) = \text{dist}(u) + w(u,v)$

$\text{parent}(v) = u$.

 endif

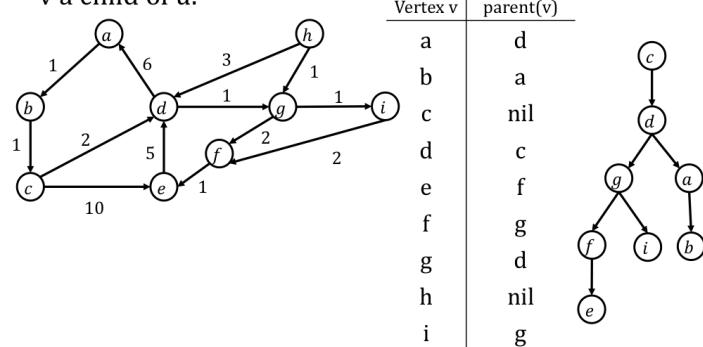
 endfor

end while

- Set $\text{parent}(v) = \text{nil}$ for all vertices $v \in V$
- Set $\text{dist}(s) = 0$ and $\text{dist}(v) = \infty$ for all other vertices $v \in V$
- Set $S = V$
- Repeat the following until S is empty
 - Remove from S the vertex u with the smallest $\text{dist}(u)$.
 /* next we relax all the outgoing edges of u */
 - For every outgoing edge (u,v) of u
 - If $\text{dist}(v) > \text{dist}(u) + w(u,v)$ then
 - Set $\text{dist}(v) = \text{dist}(u) + w(u,v)$, and $\text{parent}(v) = u$

Constructing the SP tree

- For every vertex v , if $u = \text{parent}(v)$ is not nil, the make v a child of u .



Correctness and Running Time

- It will be left as an exercise for you to prove that Dijkstra's algorithm is correct
- Just as equally instructive is an exercise for you to implement Dijkstra's algorithm in $O(|V|+|E|)\log|V|$ time. Why?
- You have already learned all the data structure for this purpose. Now it is time to practice using them.

时间复杂度: $O((|V|+|E|) * \log|V|)$

$$\sum_{u \in V} (\log|V| + d^+(u) \cdot 2\log|V|)$$

$$= \log|V| \sum_{u \in V} (1 + 2d^+(u))$$

$$= \log|V| * (|V| + 2|E|)$$

Minimum Spanning Tree (MST)

Undirected Weighted Graph

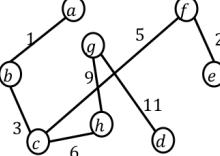
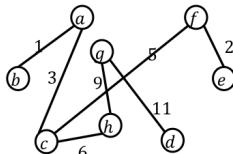
- Let $G=(V, E)$ be an undirected graph. Let w be a function that maps each edge of G to a positive integer value. Specifically, for each edge e , $w(e)$ is a positive integer value, which we call the weight of e .
- An undirected weighted graph is defined as the pair (G, w) .
- We will denote an edge between vertices u and v in G as $\{u, v\}$, instead of (u, v) , to emphasize that the ordering of u, v does not matter.
- We consider that G is connected, namely, there is a path between any two vertices in V .

Spanning Tree : Connected, undirected graph without cycles

- Given a connected undirected weighted graph (G, w) with $G=(V, E)$, a spanning tree T is a tree satisfying the following conditions:
 - The vertex set of T is V .
 - Every edge of T is an edge of G .
- The cost of T is defined as the sum of the weights of all the edges in T (note that T must have $|V|-1$ edges)

Minimum Spanning Tree

- The minimum spanning tree problem
- Given a connected undirected weighted graph (G, w) with $G=(V, E)$, the goal of the minimum spanning tree (MST) problem is to find a spanning tree of the smallest cost.
- Such a tree is called an MST of (G, w)



- Both trees are MSTs. This means that MSTs may not be unique.

给定一个连通的无向有权图 $(G=(V, E), w)$. 最小生成树就是建树时花费最小.

Prim's Algorithm

- The algorithm grows a tree T_{mst} by including one vertex at a time, at any moment, it divides the vertex set V into two parts:
 - The set S of vertices that are already in T_{mst}
 - The set of other vertices: $V \setminus S$
- at the end of the algorithm, $S = V$
- If an edge connects a vertex in S and a vertex in $V \setminus S$, we call it an extension edge.
- At all times, the algorithm enforces the following lightest extension principle:
 - For every vertex $v \in V \setminus S$, it remembers which extension edge of v has the smallest weight, referred to as the lightest extension edge of v , and denoted as $best-ext(v)$.

生成 T_{mst} 时每次加入一个新节点, 最终全部节点都加入 T_{mst} 中.

对某边, 若其一端在 S 中, 另一端在 $V \setminus S$ 中, 则该边被称作外延边.

外延边中有最轻权值的记为 $best-ext(v)$.

- Let $\{u,v\}$ be an edge with the smallest weight among all edges
- Set $S = \{u,v\}$. Initialize a tree T_{mst} with only one edge $\{u,v\}$.

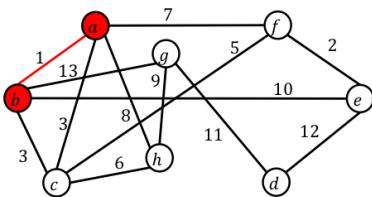
3. Enforce the lightest extension principle:

- For every vertex z of $V \setminus S$
 - If z is a neighbor of u , but not of v
 $\text{best-ext}(z) = \text{edge } \{z, u\}$
 - If z is a neighbor of v , but not of u
 $\text{best-ext}(z) = \text{edge } \{z, v\}$
 - If z is a neighbor of both u and v
 $\text{best-ext}(z) = \text{the lighter edge between } \{z, u\} \text{ and } \{z, v\}$

4. Repeat the following until $S = V$:

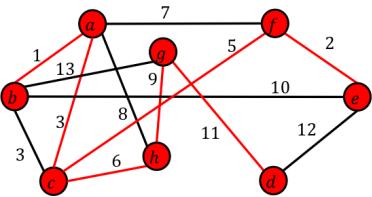
- Get an extension edge of $\{u, v\}$ with the smallest weight
/* Without loss of generality, suppose $u \in S$, and */
- Add v to S , and add edge $\{u, v\}$ into T_{mst}
/* Next, we restore the lightest extension principle. */
- For every edge $\{v, z\}$ of v :
 - If $z \notin S$ then
 - If $\text{best-ext}(z)$ is heavier than edge $\{v, z\}$ then
 - Set $\text{best-ext}(z) = \text{edge } \{v, z\}$

- Edge $\{a,b\}$ is the lightest of all. So, at the beginning $S = \{a, b\}$. The MST we are growing now has one edge $\{a,b\}$



- Note: edge $\{c,a\}$ and $\{c,b\}$ have the same weight. Either of them can be $\text{best-ext}(c)$.

- Finally, edge $\{d,g\}$ is the lightest extension edge. So, we add d to S , which now $S = \{a,b,c,f,e,h,g,d\}$, add edge $\{d,g\}$ into MST



- We have obtained our final MST.

从最轻权值边 $\{u,v\}$ 开始, 令 $S=\{u,v\}$, $T_{\text{mst}} \not\models \emptyset$
有一边 $\{u,v\}$

```

for z in V \ S
  if z in u.neighbour but not of v:
    best-ext(z) = edge(z,u)
  if z in v.neighbour but not of u:
    best-ext(z) = edge(z,v)
  if z in u.neighbour & z in v.neighbour:
    best-ext(z) = min(w(z,u), w(z,v))
  edge
while(S != V)
  得到 S 中有最小权值的外延边 (k,w).
  S.add(w); Tmst.add((k,w)).
  for every edge (w,z) of w:
    if z not in S:
      if best-ext(z) > w(z,z):
        best-ext(z) = (w,z)
    endif
  endfor
endwhile

```

Vertex v	best-ext(v) and weight
a	n/a
b	n/a
c	{c,a}, 3
d	nil, ∞
e	{e,b}, 10
f	{a,f}, 7
g	{g,b}, 13
h	{a,h}, 8

Vertex v	best-ext(v) and weight
a	n/a
b	n/a
c	n/a
d	n/a
e	n/a
f	n/a
g	n/a
h	n/a

Time Complexity Analysis

- ◆ A priority queue Q (min-heap) was employed in Prim's algorithm, what is the key of node in Q?
- ◆ Line 1 & 2: $O(1)$
- ◆ Line 3: $O(|E|)$
- ◆ Line 4: $O(|V|)$
- ◆ Line 5: $O(|V| \log |V|)$
- ◆ Line 6: $O(|V|)$
- ◆ Line 7: $O(|E| \log |V|)$, Total: $O((|V|+|E|) \log |V|)$
- ◆ Remark: Using the Fibonacci Heap, will not cover in this course, we can improve the running time to $O(|V| \log |V| + |E|)$

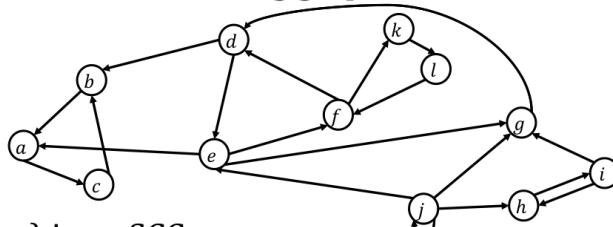
Correctness Proof

- ◆ **Claim:** For any $i \in [1, |V|-1]$, there must be an MST containing all the first i edges chosen by the algorithm
- ◆ Then the algorithm's correctness follows from the above claim at $i = |V|-1$
- ◆ We prove it by induction the sequence of the edges added to the tree
- ◆ Base case: $i=1$, let $\{u,v\}$ be the edge with the smallest weight in the graph, the edge must exist in some MST
- ◆ Inductive case: the claim holds for $i \leq k-1$
- ◆ We prove it also hold for $i=k$

Strongly Connected Components (SCC)

- ◆ Let $G=(V,E)$ be a directed graph.
- ◆ A strongly connected component (SCC) of G is a subset S of V such that:
 - ◆ For any two vertices $u, v \in S$, it must hold that:
 - ◆ There is a path from u to v
 - ◆ There is a path from v to u
 - ◆ S is maximal in the sense that we cannot put any more vertex into S without violating the above property
- ◆ It seems to be rather difficult at first glance, the algorithm is once again very simple, run DFS only twice.

- ◆ Consider the following graph:



- ◆ $\{a,b,c\}$ is an SCC
- ◆ $\{a,b,c,d\}$ is not an SCC
- ◆ $\{d,e,f,k,l\}$ is not an SCC (why?)
- ◆ $\{e,d,f,k,l,g\}$ is an SCC

给定一个有向图 $G(V,E)$.

强连通分量即是 $S \subseteq V$, S 中的每个节点之间都有路径, 且新加任何节点都会破坏这个性质.

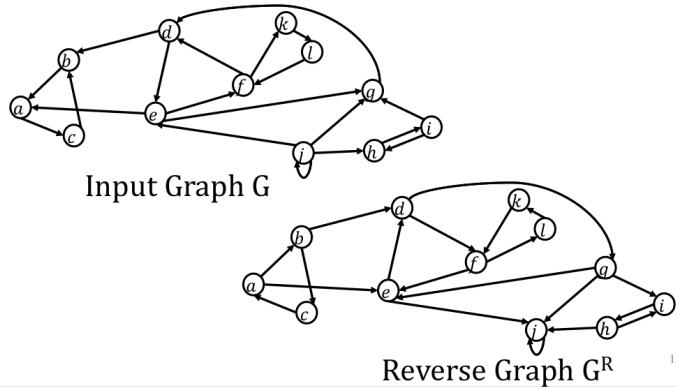
只需跑两次DFS.

SCCs are Disjoint

- Theorem: Suppose that S_1 and S_2 are both SCCs of G , Then $S_1 \cap S_2 = \emptyset$
- Proof: Assume that there is a vertex v in both S_1 and S_2 . Then, for any vertex $u_1 \in S_1$ and any vertex $u_2 \in S_2$:
 - There is a path from u_1 to u_2 : we can first go from u_1 to v within S_1 , and then from v to u_2 within S_2 .
 - Likewise, there is also a path from u_2 to u_1 .Hence, neither S_1 and S_2 is maximal, contradicting the fact that they are SCCs.

Finding SCCs Algorithm

- Step 1: obtain the reverse graph G^R by reversing the directions of all the edges in G .



- Step 2: Perform DFS on G^R , and obtain the sequence L^R that the vertices in G^R turn red (i.e., whenever a vertex is popped out of the stack, append it to L^R)
- Obtain L as the reverse order of L^R
- We may perform DFS starting from any vertex. The following is a possible order that the vertices are discovered: $f, l, k, e, j, d, g, i, h, a, b, c$
- The corresponding turn-red sequence is
 - $L^R = \{k, l, j, h, i, g, d, e, f, c, b, a\}$
 - Hence $L = \{a, b, c, f, e, d, g, i, h, j, l, k\}$
- Step 3: Perform DFS on the original graph G by obeying the following rules:
 - Rule 1: start the DFS at the first vertex of L
 - Rule 2: whenever a restart is needed, start from the first vertex of L that is still white.
- Output the vertices in each DFS-tree as an SCC

- From the last step, we have $L = \{a, b, c, f, e, d, g, i, h, j, l, k\}$

- The original graph G :
- Starting DFS from a , which discovered $\{a, b, c\}$
- Restart from f , which discovered $\{f, k, l, d, e, g\}$
- Restart from i , which discovered $\{i, h\}$
- Restart from j , which discovered $\{j\}$
- The DFS returns 4 DFS-tree, whose vertex sets are as above, Each vertex set constitutes an SCC.

Running Time Analysis

- Steps 1 and 2 obviously require only $O(|V|+|E|)$ time.
- Regarding Step 3, the DFS itself takes $O(|V|+|E|)$, but how about the cost of implement Rule 2.
- Namely, whenever, DFS needs a restart, how do we find the first white vertex in L efficiently?
- It can be done in $O(|V|)$ total time.
- Hence, the overall execution time is $O(|V|+|E|)$

Correctness Proof

- Let G be the input directed graph, with SCCs S_1, S_2, \dots, S_t for some $t \geq 1$
- Let us define a SCC graph G^{SCC} as follows:
 - Each vertex in G^{SCC} is a distinct SCC in G.
 - Consider two vertices S_i and S_j , G^{SCC} has an edge from S_i to S_j if and only if:
 - $i \neq j$
 - There is a path in G from a vertex in S_i to a vertex in S_j
- G^{SCC} is a DAG, define an SCC as a sink SCC if it has no outgoing edge in G^{SCC}
- Lemma: There must be at least one sink SCC in G^{SCC}
- Let S be a sink SCC in G^{SCC} . Suppose that we perform a DFS starting from any vertex in S. Then the first DFS-tree output must include all and only the vertex in S.
- Finding SCC: The strategy
 1. Performing DFS from any vertex in a sink SCC S
 2. Delete all vertices of S from G, as well as their edges
 3. Accordingly, delete S from G^{SCC} , as well as its edges.
 4. Repeat from Step 1, until G is empty.
- Lemma: Let S_1, S_2 be SCCs such that there is a path from S_1 to S_2 in G^{SCC} . In the ordering of L, the earliest vertex in S_2 must come before the earliest vertex in S_1