

① n nodes $\rightarrow n-1$ edges

② 若 tree 的所有 internal node 都至少有 2 个子结点, 且 leaf nodes = m ,

则 internal nodes 至少有 $m-1$ 个.

tree 中的每个 node 都是更小 tree 的 root.

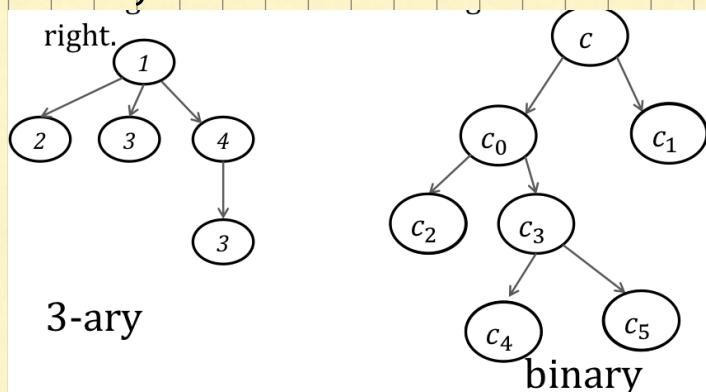
path: 从 root 到每个 node 都有唯一的 path.

depth: 到 root 的 edges 数

level: 有相同 depth 的 nodes 组成 level.

height: 最长的 depth

k -ary tree: internal tree 有着最多 k 个 child nodes 的 rooted tree.



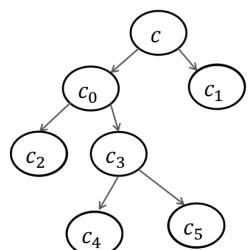
Binary tree:

(1) 空树, 或

(2) 一个 node 有一些数据, 其下有左、右子树.

Full Level: 当前 level 之上的 level 中的每个 node 都有 2 个 child node (0 level 除外)

Consider a binary tree with height h , its level l ($0 \leq l \leq h$) is full if it contains 2^l nodes.

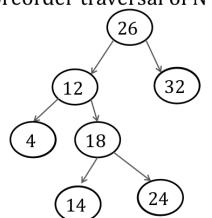


Complete Binary Tree: h height 中有 $0 \sim h-1$ level 是 full 的

③ 有 $n \geq 2$ 节点的完全二叉树有高度 $O(\log n)$.

Preorder Traversal

- Preorder traversal of the tree whose root is N
 - visit the root N
 - Recursively perform a preorder traversal of N's left subtree
 - Recursively perform a preorder traversal of N's right subtree
- Preorder traversal
 - 26, 12, 4, 18, 14, 24, 32

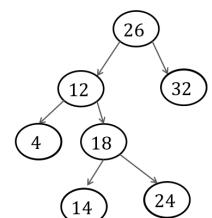


preorderiterative(treeNode root):

```
1. treeNode stack s
2. s.push(root)
3. while(s!=empty)
4.     treeNode node= s.top()
5.     print(node)
6.     s.pop()
7.     if(node->right!=null)
8.         s.push(node->right)
9.     if(node->left!=null)
10.        s.push(node->left)
```

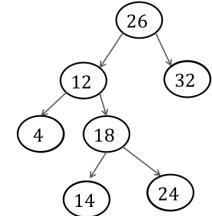
Postorder Traversal

- Postorder traversal of the tree whose root is N
 - Recursively perform a postorder traversal of N's left subtree
 - Recursively perform a postorder traversal of N's right subtree
 - visit the root N
- Postorder traversal
 - 4, 14, 24, 18, 12, 32, 26



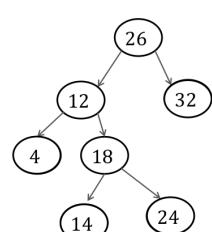
Inorder Traversal

- Inorder traversal of the tree whose root is N
 - Recursively perform a inorder traversal of N's left subtree
 - Visit the root N
 - Recursively perform a inorder traversal of N's right subtree
- Inorder traversal
 - 4, 12, 14, 18, 24, 26, 32



Level Traversal

- Visit the nodes one level at a time, from top to bottom, and left to right.
- Level-order of the tree:
 - 26, 12, 32, 4, 18, 14, 24
- How to implement?



- Preorder: root, left subtree, right subtree
- Postoder: left subtree, right subtree, root
- Inorder: left subtree, root, right subtree
- Level-order: top to bottom, left to right

Huffman Encoding

Prefix (free) code: 没有任何前缀码是另一个前缀码的前缀

- ◆ Huffman encoding uses a binary tree:

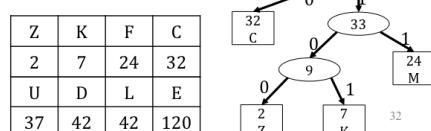
- ◆ to determine the encoding of each character
 - ◆ to decode an encoded file

- ◆ Example:

- ◆ Leaf nodes are characters

- ◆ 101 = 'D'

- ◆ "000110" = "EEEL"



- ◆ 1) Begin by reading through the text to determine the frequencies
- ◆ 2) Create a list of nodes that contain (character, frequency) pairs for each character that appears in text
- ◆ 3) Remove and "merge" the nodes with the two lowest frequencies, forming a new node that is their parent
- ◆ 4) Add the parent to the list of nodes
- ◆ 5) Repeat steps 3) and 4) until there is only a single node in the list, which will be the root of the Huffman tree.
- ◆ Example: build the Huffman tree for the following (character, frequency) pairs:

Z	K	F	C	U	D	L	E
2	7	10	12	27	30	43	65

Huffman tree 的正确性

- ◆ Given a Huffman tree, it includes at least 2 nodes, assume node u and node v have the top-2 lowest frequencies, then

- ◆ 1) node u and v have the same parent node
 - ◆ 2) $\text{depth}(u) \text{ and } \text{depth}(v) \geq \text{depth}(x)$, where node x is any leaf node in the Huffman tree.

- ◆ Huffman encoding is the optimum prefix code, i.e., the space cost is minimized.

Priority Queue =

- ◆ A **priority queue** stores a set S of n integers and supports the following operations:
 - ◆ *Insert(e)*: adds a new integer to S
 - ◆ *Delete-min*: removes the smallest integer in S, and returns it.

By Using binary heap to implement

- ◆ O(n) space consumption
- ◆ O(log n) insertion time
- ◆ O(log n) delete-min time

最小堆

- ◆ Let S be a set of n integers. A binary heap on S is a binary tree T satisfying:
 - (1) T is complete
 - (2) Every node u in T corresponds to a distinct integer in S, the integer is called the key of u (and is stored at u)
 - (3) If u is an internal node, the key of u is smaller than those of its child nodes

插入

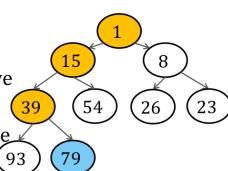
- ◆ We perform insert(*e*) on a binary heap T as follows:
 - ◆ Step 1: Create a leaf node z with key e, while ensuring that T is a complete binary tree, it means there is only one place where z could be added.
 - ◆ Step 2: Set u \leftarrow z
 - ◆ Step 3: If u is the root, return.
 - ◆ Step 4: If the key of u $>$ the key of its parent p, return
 - ◆ Step 5: Otherwise, swap the keys of u and p. Set u \leftarrow p, and repeat from Step 3.

删除

- ◆ We perform delete-min on a binary heap T as follows:
 - ◆ Step 1: Report the key of the root
 - ◆ Step 2: Identify the rightmost leaf z at the bottom level of T
 - ◆ Step 3: Delete z, and store the key of z at the root
 - ◆ Step 4: Set u \leftarrow the root
 - ◆ Step 5: If u is leaf, return
 - ◆ Step 6: If the key of u $<$ the keys of the children of u, return
 - ◆ Step 7: Otherwise, let v be the child of u with a smaller key Swap the keys of u and v. Set u \leftarrow v, and repeat from Step 5

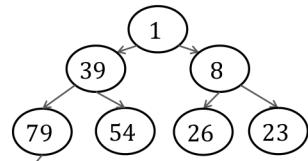
如何找到 rightmost leaf. ($O(\log n)$)

- ◆ Write the value n in binary form. We can do that in $O(\log n)$ time.
- ◆ Skip the most significant bit. We will scan the remaining bits from left to right, start from root,
 - ◆ If the bit is 0, we go to the left child of the current node
 - ◆ Otherwise, go to right child
- ◆ Here n = 9, binary form: 1001
- ◆ Skip the first bit '1'
- ◆ We scan the remaining bits
- ◆ Start from root node 1.
- ◆ The 2nd leftmost bit is 0, so we visit left, and go to node 15
- ◆ The 3rd leftmost bit is 0, so we visit left, and go to node 39
- ◆ The 4th leftmost bit is 1, so we turn right, and go to node 79 (done).



用二叉堆排序 $O(n\log n) = \text{build heap } O(n\log n) + n \times \text{delete } O(\log n)$

用数组实现完全二叉树



1	39	8	79	54	26	23	93
---	----	---	----	----	----	----	----

Storing in an array

93 Complete Binary Tree

26

- ❖ Put nodes at a higher level before those at a lower level
- ❖ Within the same level, order the nodes from left to right

假设数组下标从 1 开始

- ❖ Lemma: Suppose that node u of T is stored at $A[i]$. Then, the left child of u is stored at $A[2i]$, and the right child at $A[2i+1]$.
- ❖ Corollary: Suppose that node u of T is stored at $A[i]$. Then, the parent of u is stored at $A[\lfloor i/2 \rfloor]$.
- ❖ Lemma: the rightmost leaf node at the bottom level is stored at $A[n]$.

当整棵树都建起来之前，没必要实现 delete-min，因此，在建树时结果快。

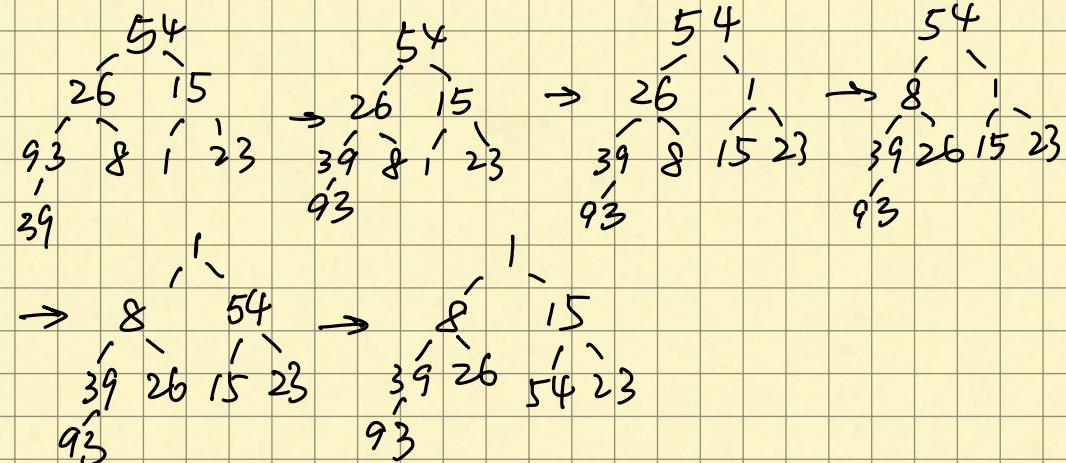
Root-fix operation

- ❖ We are given a complete binary tree T with root r . It is guaranteed that:
 - ❖ The left subtree of r is binary heap
 - ❖ The right subtree of r is a binary heap
 - ❖ However, the key of r may not be smaller than the keys of its children.
- ❖ We define the root-fix operation, it fixes the issue and makes T a binary heap.
- ❖ Root-fix can be done in $O(\log n)$ time – in the same manner as the delete-min algorithm (step 4 - 7)

54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39

Root-fix

54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39
54	26	15	93	8	1	23	39



root-fix 和 insertion 建堆是不一样的

Binary Search Tree (BST)

Balanced BST (平衡二叉查找树)

- ◆ A predecessor query: give an integer q , find its predecessor in S , which is the largest integer in S that does not exceed q .
- ◆ Insertion: adds a new integer to S
- ◆ Deletion: removes a integer from S

BST 中的复杂度

- ◆ $O(n)$ space consumption
- ◆ $O(h)$ time per predecessor query (hence, also per dictionary lookup)
- ◆ $O(h)$ time per insertion
- ◆ $O(h)$ time per deletion
- ◆ where $n = |S|$, h is the height of BST, Note that all the above complexities hold in the worst case.

BST

- ◆ A BST on a set S of n integers in a binary tree T satisfying all the following requirements:
 - ◆ T has n nodes
 - ◆ Each node u in T stores a distinct integer in S , which is called the key of u
 - ◆ For every internal u , it holds that:
 - ◆ The key of u is larger than all the keys in the left subtree of u .
 - ◆ The key of u is smaller than all the keys in the right subtree of u .

BST 不一定是完全二叉树

Predecessor Query

- ◆ Suppose that we have created a BST T on a set S of n integers. A predecessor query with search value q can be answered by descending a single root-to-leaf path:
 - ◆ (1) Set $p \leftarrow -\infty$ (p will contain the final answer at the end)
 - ◆ (2) Set $u \leftarrow$ the root of T
 - ◆ (3) If $u = \text{nil}$, then return p
 - ◆ (4) If key of $u = q$, then set p to q , and return p
 - ◆ (5) If key of $u > q$, then set u to the left child (now $u = \text{nil}$ if there is no left child), and repeat from Step (3)
 - ◆ (6) Otherwise, set p to the key of u and u to the right child, and repeat from Step (3)

每一次访问节点都是 $O(1)$
故若 BST 深 h 时，最坏 $O(h)$.

找小于 q 的最大值

Successor Query

- ◆ The successors of n integer q in S is the smallest integer in S that is no smaller than q .

找大于 q 的最小值

5 Predecessor Query 类似。

BST Insertion

- Suppose that we need to insert a new integer e. First create a new leaf z storing the key e. This can be done by descending a root-to-leaf path:
 - 1. Set u \leftarrow the root of T
 - 2. If $e <$ the key of u
 - 2.1 If u has a left child, then set u to the left child
 - 2.2 Otherwise, make z the left child of u, and done
 - 3. Otherwise:
 - 3.1 If u has a right child, then set u to the right child
 - 3.2 Otherwise, make z the right child of u, and done.
 - Repeat from Step 2.
- The total cost is proportional to the height of T, i.e., $O(h)$

BST Deletion

- Suppose that we want to delete an integer e. First, find the node u whose key equals to e in $O(h)$ time (through a predecessor query).
- Case 1: if u is a leaf node, simply remove it from T.
- Case 2: if u has a right subtree:
 - Find the node v storing the successor s of e.
 - Set the key of u to s
 - Case 2.1: if v is a leaf node, then remove it from T
 - Case 2.2: otherwise, it must hold that v has a right child w, but not left child. Replace node v by subtree which rooted at w.
- Case 3: if u has no right subtree:
 - It must hold that u has a left child v. Replace node u by the subtree rooted at v.
- In all above cases, we have essentially descended a root-to-leaf path (call it deletion path), and removed a leaf node.
- The cost so far is $O(h)$, recall that the successor of an integer can be found in $O(h)$ time.

BBST

- A binary tree T is balanced if the following holds on every internal node u of T:
 - The height of the left subtree of u differs from that the right subtree of u by at most 1.
- Theorem: a balanced binary tree with n nodes has height $O(\log n)$.

AVL-tree

- An AVL-tree on a set S of n integers is a balanced binary search tree T, where the following hold on every internal node u
 - u stores the heights of its left and right subtrees.

Red-black tree

Splay tree

AVL tree

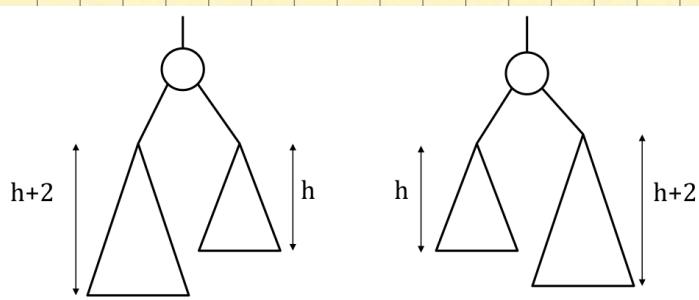
先通过 predecessor / successor query 找到要删除的结点。

若该结点是叶子，就直接删除。

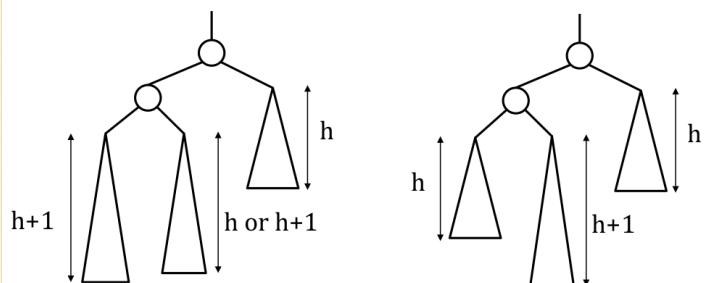
若该结点为内结点，由 predecessor / successor 的性质可知，若有右子树/左子树，则必无左子树/右子树

$$f_{\text{root}, n} = \max(f_{\text{left}, n-1}, f_{\text{right}, n-1}) + 1$$

2-level imbalance



仅讨论左子树高2的情况，右子树高2对称。

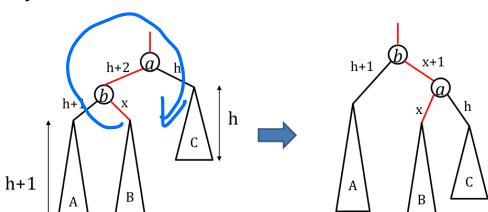


Left-Left case

Left-Right case

Rebalance Left-Left

- By a rotation:

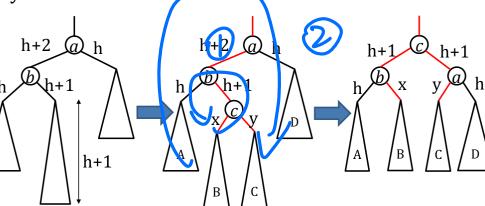


- Only 3 pointers to change (the red ones). The cost is $O(1)$.

- Recall that $x = h$ or $h+1$

Rebalance Left-Right

- By a double rotation:



- Only 5 pointers to change (see above). Hence, the cost is $O(1)$.

- Note that x and y must be h or $h-1$. Furthermore, at least one of them must be h (why?)

Insertion time analysis

- It will be left as an exercise for you to prove
 - Only 2-level imbalance can occur in an insertion
 - Once we have remedied the lowest imbalance node, all the nodes in the tree will become balanced again
- Thus, we can conclude the insertion cost in a balanced BST is $O(\log n)$, why?

Deletion time analysis

- It will be left as an exercise for you to prove
 - Only 2-level imbalance can occur after a deletion
- Thus, we can conclude the deletion cost in a balanced BST is $O(\log n)$

- $O(n)$ space consumption

- $O(\log n)$ time per predecessor query (hence, also per dictionary lookup)

- $O(\log n)$ time per insertion

- $O(\log n)$ time per deletion